

## **Paper Review: Implicit Quantile Networks for Distributional Reinforcement Learning**

### **Summary:**

The paper "Implicit Quantile Networks for Distributional Reinforcement Learning" introduces a new approach for reinforcement learning (RL) that learns the distribution of returns rather than the expected value. The method, called Implicit Quantile Networks (IQN), estimates the quantile function of the return distribution using a deep neural network. IQN outperforms previous methods in the same area of RL. The paper covers several benchmark tasks in RL and provides analysis and evaluation of the experimental results found in the figures.

### **Contributions:**

The paper makes two main contributions to the field of reinforcement learning (RL). First, it proposes a new approach to distributional RL that estimates the quantile function of the return distribution, rather than the expected value; this adds value to their contributions. Second, the paper demonstrates the superior performance of IQN compared to previous methods in several benchmark tasks in RL. Overall, both contributions open new possibilities for research in distributional RL and provide a promising direction for developing more effective RL and ML algorithms.

### **Strengths and Weaknesses:**

The strengths of the paper include its clear explanation and breakdown of the new approach, its in-depth evaluation of the method, and its empirical demonstration of IQN's superior performance on several benchmark tasks. The paper's weaknesses lie in its' lack of a detailed analysis of the method's constraints and the absence of a comparison with other recent approaches to distributional RL; this is crucial as it is not comparing itself to potential algorithms in its same level of complexity.

### **Experimental Validity:**

The paper's experimental validation is compelling. The authors evaluate IQN on several benchmark tasks and demonstrate its superior performance compared to previous state-of-the-art methods. They also perform a testbench analysis to show the importance of various components of the IQN architecture. This alone validates the experiment and allows for a clear replication of the experiment for others to conduct.

### **How can this work be extended:**

The paper opens several directions for future research in RL. One potential extension is to explore the use of IQN in other types of RL problems, such as continuous control tasks. Another potential extension is to investigate the use of IQN in multi-agent RL settings. Finally, it would be

interesting to explore IQN in more complicated video games like 3D rendered ones today, that could allow for multiple agents.