



The Bradley Department

Electrical & Computer Engineering

ECE 5984 – Deep Reinforcement Learning

Project Assignment #2: SARSA & Q-learning on FrozenLake

SAR and Q-learning:

Both SARA and Q-learning algorithms are value-based temporal difference methods, which solve the problem with iteration over the full set of the state-action pairs. We use states and actions obtained from the environment to update the values of state-action pairs, i.e., q values. In this assignment, you are asked to implement the SARSA & Q-learning algorithms to solve the frozen lake problem.

I. SARSA:

A. Set up/import Python libraries for this project

```
import gym
import collections

ENV_NAME = "FrozenLake-v1"
GAMMA = 0.9
ALPHA = 0.2
TEST_EPSOIDS = 200
```

B. Implement an Agent_SARSA class outlined below:

```
class Agent_SARSA:
    def __init__(self):
        self.env = gym.make(ENV_NAME)
        self.state = self.env.reset()
        self.values = collections.defaultdict(float)

    def sample_env(self):
        """
        Inputs:
            - self: an agent

        Returns:
```

```

        - a tuple: (old_state, action, reward, new_reward)
    """

def choose_action(self, state):
    """
    Inputs:
        - self: an agent
        - state: current state

    Returns:
        - next_a: the next action taken.
    """

def value_update(self, s, a, r, next_s, next_a):
    """
    Inputs:
        - self: an agent
        - s: state
        - a: action
        - r: reward
        - next_s: next state

    Returns:
        - self.values[(s, a)]: the updated value of (s, a).
    """

def play_episode(self, env):
    """
    Inputs:
        - self: an agent
        - env: the environment

    Returns:
        - total_reward: the total reward after playing an
        episode
    """

```

- C. Implement the train loop: we first create a test environment and agent, then in the loop, we do one step in the environment and perform a value update using the obtained data. We test the current policy by playing several test episodes. If a good reward is obtained, then we stop the training.

II. Q-learning:

- A. Set up/import Python libraries for this project

```
import gym
import collections

ENV_NAME = "FrozenLake-v0"
GAMMA = 0.9
ALPHA = 0.2
TEST_EPSOIDS = 200
```

- B. Implement an Agent_QLearning class outlined below:

```
class Agent_QLearning:
    def __init__(self):
        self.env = gym.make(ENV_NAME)
        self.state = self.env.reset()
        self.values = collections.defaultdict(float)

    def sample_env(self):
        """
        Inputs:
            - self: an agent

        Returns:
            - a tuple: (old_state, action, reward, new_reward)
        """

    def best_value_and_action(self, state):
        """
        Inputs:
            - self: an agent
            - state: current state

        Returns:
            - best_value: the best value updated
        """
```

```

        - best_action: the best action taken.
    """

def value_update(self, s, a, r, next_s):
    """
    Inputs:
        - self: an agent
        - s: state
        - a: action
        - r: reward
        - next_s: next state

    Returns:
        - self.values[(s, a)]: the updated value of (s, a).
    """

def play_episode(self, env):
    """
    Inputs:
        - self: an agent
        - env: the environment

    Returns:
        - total_reward: the total reward after playing an
          episode
    """

```

C. Implement the train loop: we first create a test environment and agent, then in the loop, we do one step in the environment and perform a value update using the obtained data. We test the current policy by playing several test episodes. If a good reward is obtained, then we stop the training.

You need to prepare **a written report** (in the pdf format) including the following sections: (1) Approaches, (2) Experimental Results, and (3) Discussion.

In addition, you need to attach your **implementation codes/notebooks** as separate files to the report.